Econometrics Math Camp
Day One: Foundations of Statistics

Michael Droste

August 21, 2020

Introduction

- The goal of the next two days is to get you acquainted with some of the tools you will be using throughout the first-year econometrics sequence.

- The materials for the econometrics math camp (slides, notes, and problems) are available on GitHub: https://github.com/mdroste/metrics-mathcamp-2020

- These notes and slides were developed by Ashesh Rambachan, Frank Pinter, and myself. All errors are mine - let me know if you find any!

- These slides will largely consist of 'highlights' from the econometrics handbook; that is, the concepts that I think will be most useful for you to know before starting the first year sequence.

- Please have both a computer and a paper/tablet ready!

Today's Outline

- Probability
    - Random Experiments
    - Probabilities and Conditional Probabilities

- Random Variables
    - Definition
    - Continuous and Discrete Random Variables
    - Cumulative Distribution Functions
    - Joint and Marginal Distributions
    - Conditioning and Independence
    - Transformations of Random Variables

- Expectations
    - Definition and Properties
    - Conditional Expectations

Motivation

- We will develop a theory of probability by using some tools from a branch of mathematics called measure theory, which you may not have seen before. Measure theory provides a unified lens through which we can think about probability, and it requires very little mathematical structure.

- After we have defined notions of probability, random variables, etc. in measure-theoretic terms, we can largely forget about this abstract machinery, in the sense that virtually all of your first-year homework will consist of basic algebra and applying laws of probability.

Random Experiments

- We are interested in a very general concept that we will call a random experiment: an experiment or process whose outcome is not known to us beforehand.

- The two fundamental building blocks of the random experiment are the sample space of the experiment and the events of the experiment.

- Let Ω denote the sample space of the random experiment. The sample space Ω is the set of all possible outcomes of the experiment we are studying.

- We will call a subset $A \subseteq \Omega$ an event of the random experiment. It will also be helpful to let $\mathcal{A}$ denote the set of all events (i.e., $\mathcal{A}$ is the set containing all subsets of Ω).

Random Experiments: Examples

- To make these definitions more salient, let's consider two quick examples.

- Example 1: Suppose we survey 10 randomly selected people on their employment status and count how many are unemployed - that is, our outcome is the count of people who are unemployed.

    1. What is the sample space Ω?

    2. What is the event A such that more than 30% of those surveyed are unemployed?

Random Experiments: Examples

- To make these definitions more salient, let's consider two quick examples.

- Example 1: Suppose we survey 10 randomly selected people on their employment status and count how many are unemployed.

  1. What is the sample space Ω?
     $\Omega = \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10\}$
  2. What is the event A such that more than 30% of those surveyed are unemployed?
     $A = \{4, 5, 6, 7, 8, 9, 10\}$

Random Experiments: Examples

- To make these definitions more salient, let's consider two quick examples.

- Example 1: Suppose we survey 10 randomly selected people on their employment status and count how many are unemployed.

  1. What is the sample space Ω?
     $\Omega = \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10\}$
  2. What is the event A such that more than 30% of those surveyed are unemployed?
     $A = \{4, 5, 6, 7, 8, 9, 10\}$

- Example 2: Suppose I survey a random person on their income (Raj was unavailable).

  1. What is the sample space Ω?

  2. How would I write down the event A such that a person earns between $30k and $40k?

Random Experiments: Examples

- To make these definitions more salient, let's consider two quick examples.

- Example 1: Suppose we survey 10 randomly selected people on their employment status and count how many are unemployed.

  1. What is the sample space Ω?
     $\Omega = \{0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10\}$
  2. What is the event A such that more than 30% of those surveyed are unemployed?
     $A = \{4, 5, 6, 7, 8, 9, 10\}$

- Example 2: Suppose I survey a random person on their income (Raj was unavailable).

  1. What is the sample space Ω?
     $\Omega = \mathbb{R}$
  2. How would I write down the event A such that a person earns between $30k and $40k?
     $A = [\$30000, \$40000]$

Probability

- Now that we have developed some definitions to characterize random experiments, our goal is to sensibly define the probability of an event.

- Before we begin - when we refer to the probability of an event, what do you think we are trying to communicate?

- Why is this not trivial? If the sample space Ω is finite (example #1), probability is easy to define. But when the sample space is not finite - for instance, if a person's income can take on any real number (see example #2) - then we need to think more carefully.

- It turns out that measure theory is the appropriate mathematical framework to define probabilities in a unified way. Frankly, you will not see much measure theory (or any) in the first-year sequence.

Probability: Defining σ-algebras

- Let Ω be a set and $\mathcal{A} \subseteq 2^{\Omega}$ be a family of its subsets. We say that $\mathcal{A}$ is a σ-algebra if (and only if) it satisfies:

    1. $\Omega \in \mathcal{A}$
    2. Closure under complements: If $A \in \mathcal{A}$, then $A^c = \Omega \setminus A \in \mathcal{A}$.
    3. Closure under countable unions: If $A_n \in \mathcal{A}$ for $n = 1, 2, ...$, then $\cup_n A_n \in \mathcal{A}$

- Note that we will adopt a little bit of set theory jargon. Let X^c denote the complement of a set X, and let $\cup$ and $\cap$ denote the union and intersection operators, respectively.

- If Ω is a set and $\mathcal{A} \subseteq 2^{\Omega}$ is a σ-algebra, we say that $(\Omega, \mathcal{A})$ is a measurable space.

- If $(\Omega, \mathcal{A})$ is a measurable space, we say $A \in \mathcal{A}$ is measurable with respect to $\mathcal{A}$.

Probability: Properties of σ-algebras

- σ-algebras have lots of useful properties. Here are two immediately useful ones:

    - $\emptyset \in \mathcal{A}$

    - Closure under countable intersections: If $A_n \in \mathcal{A}$ for $n = 1, 2, ...$, then $\cap_n A_n \in \mathcal{A}$.

- You will prove these in our first breakout session shortly.

Probability: Measures

- Let $(\Omega, \mathcal{A})$ be a measurable space. A measure is a function, $\mu : \mathcal{A} \to \mathbb{R}$, such that:

  1. $\mu(\emptyset) = 0$
  2. $\mu(A) \geq 0$ for all $A \in \mathcal{A}$
  3. If $A_n \in \mathcal{A}$ for $n = 1, 2, \ldots$ with $A_i \cap A_j = \emptyset$ for $i \neq j$, then $\mu(\cup_n A_n) = \sum_n \mu(A_n)$

- If $\mu(\Omega) = 1$, we say that μ is a probability measure, denoted $P : A \to [0, 1]$.

Probability: Probability Spaces

- Congrats! We have now developed all of the building blocks we need to characterize the probability of any random experiment.

- A random experiment is characterized by a probability space $(\Omega, \mathcal{A}, P)$, where:

    - Ω: The sample space, or set of outcomes

    - $\mathcal{A}$: The set of events, assumed to admit a σ-algebra representation

    - P: A probability measure defined on the σ-algebra

Probability and Measure Theory

- Why did we need to go through such abstract machinery to think about the probability of events? It turns out that all of the usual rules of probability 'pop out' of the properties of the mathematical structure in our probability space (i.e. σ-algebras, measures) almost immediately.

- In addition, your intuitive understanding of a probability (i.e. a long-run average or subjective beliefs about a process) can be shown to nest within this setup.

- We will not discuss measure theory much more in math camp or in the econometrics sequence (at other programs, they will spend more time on this). Having some understanding that probability is built on the fundamentals of measure theory is a powerful idea you should remember.

Basic Laws of Probability

- Consider a probability space $(\Omega, \mathcal{A}, P)$. The following laws always hold:

    - For any $A \in \mathcal{A}$, we have $P(A^c) = 1 - P(A)$

    - $P(\Omega) = 1$

    - If $A_1, A_2 \in \mathcal{A}$ with $A_1 \subseteq A_2$, then $P(A_1) \leq P(A_2)$

    - For all $A \in \mathcal{A}$, $0 \leq P(A) \leq P(\Omega)$

    - If $A_1, A_2 \in \mathcal{A}$, then $P(A_1 \cup A_2) = P(A_1) + P(A_2) - P(A_1 \cap A_2)$

- You will prove some of these shortly in a breakout session.

Conditional Probability

- Given a random experiment and the information that event B has occurred, what is the probability that the outcome also belongs to event A?

- Let $A, B \in \mathcal{A}$ with $P(B) > 0$. The conditional probability of A given B is:

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

- Note that all of our usual basic laws of probability apply to $P(A|B)$ - and so do the laws of algebra!

- We will see later on that conditional probabilities are really important, but they also motivate three really important rules that come up when manipulating conditional probabilities.

Conditional Probability: Multiplication Rule

- Consider n events $A_1, ..., A_n$. The multiplication rule relates the probability of all events A_i occurring jointly to conditional probabilities. In general, we can express this as:

$$P(\cap_{i=1}^n A_i) = P(A_1)P(A_2|A_1)P(A_3|A_2 \cap A_1) \cdots P(A_n| \cap_{i=1}^{n-1} A_i)$$

  - When n=2, we have $P(A_1 \cap A_2) = P(A_1)P(A_2|A_1)$. Note this is basically just rearranging the definition of a conditional probability. This is the setting in which the multiplication rule usually comes up.

  - When n=3, we have $P(A_1 \cap A_2 \cap A_3) = P(A_1)P(A_2|A_1)P(A_3|A_1 \cap A_2)$.

Conditional Probability: Law of Total Probability

- Consider K disjoint events C_k that partition the sample space Ω; that is, $C_i \cap C_j = \emptyset$ for all $i \neq j$ and $\cup_{i=1}^K C_i = \Omega$. Let A be some event.

- The law of total probability states that we can write $P(A)$ in terms of $P(A|C_i)$ and $P(C_i)$ in a way that 'adds up'.

$$P(A) = \sum_{i=1}^K P(A|C_i)P(C_i)$$

Conditional Probability: Bayes Rule

- Given two events A, B, Bayes' rule (sometimes seen as Bayes' law) relates the conditional probabilities $P(A|B)$, $P(B|A)$ and the marginal probabilities $P(A)$, $P(B)$.

- One simple formulation of Bayes law can be expressed as:

$$P(B|A) = \frac{P(A|B)P(B)}{P(A)}$$

- Proof? (hint: use the multiplication rule)

- You can write entire papers that are basically just Bayes' law...

Identification and Estimation of Undetected COVID-19 Cases Using Testing Data from Iceland

Karl M. Aspelund, Michael C. Droste, James H. Stock, Christopher D. Walker

NBER Working Paper No. 27528
Issued in July 2020
NBER Program(s): Economic Fluctuations and Growth, Health Care, Health Economics, Technical Working Papers

Independence

- What if an event B has no information about an event A?

- We say that two events A and B are independent if $P(A|B) = P(A)$, or equivalently, $P(B|A) = P(B)$, or $P(A \cap B) = P(A)P(B)$.

- Let $E_1, ..., E_n$ be events. $E_1, ..., E_n$ are said to be jointly independent if for any $i_1, ..., i_k$:

$$P(E_{i_1} | E_{i_2} \cap \cdots \cap E_{i_k}) = E_{i_1}$$

- Given an event C, we say that two events A and B are conditionally independent if:

$$P(A \cap B | C) = P(A|C)P(B|C)$$

Break-Out Session #1

- We will now split out into breakout rooms to work through proving some of the laws of probability using the mathematical machinery we built up.

- Please see the attached Metrics Math Camp Worksheet #1 for a brief primer on the definitions you'll need to prove (i.e. the definition of a probability measure and a σ-algebra) and the problems.

- Take 15 minutes to work through these problems with your small group. Raise your hand on Zoom if you get stuck or have a question.

- The actual process of proving these results is way less important than taking this as an opportunity to work with some new folks in your class.

Break

- This concludes our brief primer on the basics of probability. So far, we have:
    1. Developed a measure-theoretic model of probability
    2. Derived rules for manipulating probabilities from the structure of our model

- In the next part of math camp, we'll discuss random variables, expectations, and conditional expectations.

- Before we continue, let's take a short ($\approx$10 min) break.

- I'll be here to chat and answer any questions!

Random Variables

- You're going to hear the term 'random variable' quite a bit in the first year.

- What is a random variable, though? You probably have a pretty good idea.

- Consider the random experiment we thought about earlier: a random variable might be thought as representing the outcomes of a random experiment.

- To be more precise, we're going to need two additional building blocks from measure theory.

Random Variables: Building Blocks

- The first building block we need is a particular σ-algebra, the Borel σ-algebra, often denoted $\mathcal{B}$.

- $\Omega = \mathbb{R}$, $\mathcal{A} = $ collection of all open intervals in $\mathbb{R}$. The "smallest" σ-algebra containing all open sets is the Borel σ-algebra.

- More rigorously, $\mathcal{B}$ is the collection of all Borel sets, which is any set in $\mathbb{R}$ formed by countable union, countable intersection, relative complement.

Random Variables: Building Blocks

- The second building block we need is the idea of a measurable function between two measure spaces.

- Let $(\Omega, \mathcal{A}, \mu)$ and $(\Omega', \mathcal{A}', \mu')$ two measure spaces. Let $f : \Omega \to \Omega'$ be a function. We say that f is measurable if (and only if) $f^{-1}(A') \in \mathcal{A}$ for all $A' \in \mathcal{A}'$.

- What does this even mean? Loosely, a measurable function can be thought of as a function between the sets underneath two measure spaces that preserves the structure of the measure spaces: the preimage of any measurable set is measurable.

Random Variables: Definition

- Let $(\Omega, \mathcal{A}, P)$ denote a probability space and let $X : \Omega \to \mathbb{R}$ denote a real-valued function. We say that X is a random variable if (and only if) X is P-measurable. That is, $X^{-1}(B) \in \mathcal{A}$ for all $B \in \mathcal{B}$ where $\mathcal{B}$ is the Borel σ-algebra.

Cumulative Distribution Function

- Let X be a random variable. The cumulative distribution function (or cdf) or X, $F : \mathbb{R} \to [0, 1]$, is defined as:

$$F_X(x) = P(X^{-1}(x)) = P(\{\omega \in \Omega : X(\omega) \le x\})$$

We often write the cdf of X as:

$$F_x(x) = P(X \le x)$$

- Loosely, the cdf of X is a function that tells you, for any given value of x, the probability that the random variable X takes on a value less than or equal to x.

Cumulative Distribution Function: Properties

- The cumulative distribution function F_X has many handy properties. Among them:

    1. For $x_1 \leq x_2$, $F_X(x_2) - F_X(x_1) = P(x_1 < X < x_2)$
    2. $\lim_{x \to -\infty} F_X(x) = 0$ and $\lim_{x \to \infty}(x) = 1$
    3. F_X is non-decreasing.
    4. F_X is right-continuous.

Cumulative Distribution Functions: Quantiles

- The quantiles of a random variable X are given byt he inverse of its cumulative distribution function.

- The quantile function is:
$$Q(u) = \inf\{x : F_X(x) \geq u\}$$

  If F_X is invertible, then:
$$Q(u) = F_X^{-1}(u)$$

- Cool fact: for any function F that satisfies the handy properties on the previous slide, we can construct a random variable whose cumulative distribution function is F. The details of this construction are left as an exercise in the next break-out session.

Discrete Random Variables

- Let X be a random variable. We say that X is a discrete random variable if (and only if) F_X is constant except at a countable number of points (i.e. F_X is a step function).

$$p_i = P(X = x_i) = F_X(x_i) - \lim_{x \to x_i^-} F_X(x)$$

Use this to define the probability mass function of X:

$$f_X(x) = \begin{cases} p_i & \text{if } x = x_i, i = 1, 2, \ldots \\ 0 & \text{otherwise} \end{cases}$$

- We can write:

$$P(x_1 < X \le x_2) = \sum_{x_1 < x \le x_2} f_X(x)$$

Continuous Random Variables

- Let X be a random variable. We say that X is a continuous random variable if (and only if) F_X can be written as:

$$F_X(x) = \int_{-\infty}^{\infty} f_X(t)\,dt$$

where f_X satisfies $f_X(x) \geq 0$ and $\int_{-\infty}^{\infty} f_X(t)\,dt = 1$.

- At the points where F_X is continuous, we have:

$$f_X(x) = \frac{dF_x(x)}{dx}$$

- We call $f_X(x)$ the probability density function (or pdf) of X.

- The support of X is $S_X = \{x : f_X(x) > 0\}$

Continuous Random Variables: Notes

- There are two useful facts to remember about continuous random variables (actually, probably more).

- First, note that for $x_2 \geq x_1$, we have:

$$P(x_1 < X \leq x_2) = F_X(x_2) - F_X(x_1)$$
$$= \int_{x_1}^{x_2} f_X(t)\,dt$$

- Second, note that $P(X = x) = 0$; that is, the probability that a continuous random variable takes on any particular value x is 0. At a deep level, this fact is why we needed measure theory to describe random variables!

Continuous Random Variables

- Let X be a random variable. We say that X is a continuous random variable if (and only if) F_X can be written as:

$$F_X(x) = \int_{-\infty}^{\infty} f_X(t)dt$$

where f_X satisfies $f_X(x) \geq 0$ and $\int_{-\infty}^{\infty} f_X(t)dt = 1$.

- At the points where f_X is continuous, we have:

$$f_X(x) = \frac{dF_x(x)}{dx}$$

- We call $f_X(x)$ the probability density function (or pdf) of X.

- The support of X is $S_X = \{x : f_X(x) > 0\}$

Break

- This seems like a nice time for a quick break!

- Let's take another 10 minute break.

- As before, I'm here to chat with anyone about anything.

Joint Distributions

- Let X and Y be two (scalar) random variables. A random vector (X, Y) is a measurable mapping from Ω to $\mathbb{R}^2$.

- The joint cumulative distribution function of (X, Y) is:

$$
\begin{aligned}
F_{X,Y}(x, y) =& P(X \leq x, Y \leq y) \\
=& P(\{\omega : X(\omega) \leq x\} \cap \{\omega : Y(\omega) \leq y\})
\end{aligned}
$$

We say that (X, Y) is a discrete random vector if:

$$
F_{X,Y}(x, y) = \sum_{u \leq x} \sum_{v \leq y} f_{X,Y}(u, v)
$$

where $f_{X,Y}(x, y) = P(X = x, Y = y)$ is the joint probability mass function of (X, Y).

Joint Distributions

- Let X and Y be two (scalar) random variables. We say that (X, Y) is a continuous random vector if:

$$F_{X,Y}(x, y) = \int_{-\infty}^{x} \int_{-\infty}^{y} f_{X,Y}(u, v) dv du$$

where $f_{X,Y}(x, y)$ is the joint probability density function of (X, Y).

- As before, at points where $F_{X,Y}$ is continuous, we have:

$$f_{X,Y}(x, y) = \frac{\partial^2 F_{X,Y}(x, y)}{\partial x \partial y}$$

Joint Distributions to Marginal Distributions

- If we know the joint cdf of (X, Y), we can recover the marginal *cdfs* of X and Y:

$$F_X(x) = P(X \leq x)$$
$$= P(X \leq x, Y \leq \infty)$$
$$= \lim_{y \to \infty} F_{X,Y}(x, y)$$

- We can also recover the marginal pdfs from the joint pdf:

$$f_X(x) = \sum_y f_{X,Y}(x, y) \qquad \text{(discrete)}$$
$$f_X(x) = \int_{S_y} f_{X,Y}(x, y) dy \qquad \text{(continuous)}$$

Conditioning with Discrete Variables

- Consider x with $f_X(x) > 0$. The conditional pmf of Y given $X = x$ is:

$$f_{Y|X}(y|x) = \frac{f_{X,Y}(x,y)}{f_X(x)}$$

where $f_{Y|X}(y|x)$ satisfies $f_{Y|X}(y|x) \geq 0$ and $\sum_y f_{Y|X}(y|x) = 1$.

- The conditional cdf of Y given $X = x$ is defined:

$$F_{Y|X}(y|x) = P(Y \leq y|X = x) = \sum_{v \leq y} f_{Y|X}(v|x)$$

Conditioning with Continuous Variables

- Consider x with $f_X(x) > 0$. The conditional pdf of Y given $X = x$ is:

$$f_{Y|X}(y|x) = \frac{f_{X,Y}(x, y)}{f_X(x)}$$

- The conditional cdf of Y given $X = x$ is defined:

$$F_{Y|X}(y|x) = \int_{-\infty}^{y} f_{Y|X}(v|x)\,dv$$

Independence of Random Variables

- Let X and Y be random variables. We say that X and Y are independent if (and only if):

$$F_{Y|X}(y|x) = F_Y(y)$$

  or, equivalently:

$$F_{X,Y}(x,y) = F_X(x)F_Y(y)$$

- This can also be defined in terms of densities - replace F with f above.

Transformations of Random Variables

- Let X be a random variable with a cdf F_X. Define a random variable $Y = h(x)$, where h is a one-to-one function whose inverse h^{-1} exists. What is the distribution of Y?

- First, let's tackle the discrete case. Suppose that X is discrete with values $x_1, ..., x_n$. Y is also discrete with the values $y_i = h(x_i)$ for $i = 1, ..., n$, and the pmf of Y is given by:

$$P(Y = y_i) = P(X = h^{-1}(x_i))$$
$$f_Y(y) = f_X(h^{-1}(y_i))$$

Transformations of Random Variables

- Next, let's consider the case where X is a continuous random variable. It helps to consider two cases:

    1. First, suppose h is increasing. Then we have:

$$F_Y(y) = P(Y \le y) = P(X \le h^{-1}(y)) = F_X(h^{-1}(y))$$

    and so $f_Y(y) = \frac{dF_Y(y)}{dy} = f_X(h^{-1}(y))\frac{dh^{-1}(y)}{dy}$.

    2. Second, suppose h is decreasing. Then we have:

$$f_Y(y) = -f_X(h^{-1}(y))\frac{dh^{-1}(y)}{dy}$$

- Combining these cases, we have that in general:

$$f_Y(y) = f_X(h^{-1}(y))|\frac{dh^{-1}(y)}{dy}|$$

Expectations of Discrete Random Variables

- Let X be a discrete random variable. Its expectation (or expected value) is defined as:

$$E[X] = \sum_x x f_X(x)$$

  if $\sum_x |x| f_X(x) < \infty$. Otherwise, the expectation does not exist.

- Note that expectations play nicely with transformations of random variables. For instance, let $g : \mathbb{R} \to \mathbb{R}$. Then:

$$E[g(X)] = \sum_x g(x) f_X(x)$$

Expectations of Continuous Random Variables

- Let X be a continuous random variable. Its expectation (or expected value) is defined as:

$$E[X] = \int_{S_x} x f_X(x) dx$$

if $\int_{S_x} |x| f_X(x) dx < \infty$. Otherwise, the expectation does not exist.

- Note that expectations (still) play nicely with transformations of (continuous) random variables. For instance, let $g : \mathbb{R} \to \mathbb{R}$. Then:

$$E[g(X)] = \int_{S_x} g(x) f_X(x) dx$$

Expectations as Linear Operators

- Expectations are a linear operator. What does this mean? Let X be a random variable, $a \in \mathbb{R}$ a constant, and $g_1(\cdot), g_2(\cdot)$ be real-valued functions. Then:

  1. $E[a] = a$
  2. $E[ag_1(X)] = aE[g_1(X)]$
  3. $E[g_1(X) + g_2(X)] = E[g_1(X)] + E[g_2(X)]$

Conditional Expectations

- Let X and Y be random variables with a joint density $f_{X,Y}(x,y)$. The conditional expectation of Y given X=x is:

$$E[Y|X = x] = \int_{S_Y} y f_{Y|X}(y|x) dy$$

- Note that this is a function of x, and is sometimes called the conditional expectation function or regression function. It is sometimes useful to denote the CEF of a variable Y as a function of x as $\mu_Y(x)$.

Conditional Expectations: Properties

- Conditional expectations are going to show up again and again in this course. They are intimately related to regression analysis.

- One incredibly useful property of the CEF is called the CEF decomposition. Let Y be a random variable. We can write:

$$Y = E[Y|X] + \epsilon$$

where ϵ is mean independent of X and is therefore uncorrelated with *any* function of X.

- You will try to prove this property in the second breakout session.

Conditional Expectations as Optimal Forecasts

- Conditional expectation functions are incredibly useful objects, and they have several useful interpretations in econometrics.

- One interpretation is that conditional expectations are the solution to an optimal forecasting problem. Suppose you want to forecast the value of a random variable Y. More precisely, suppose you want to pick some $h \in \mathbb{R}$ that minimizes the expected mean squared error:

$$E[(Y - h)^2] = \int (y - h)^2 f_Y(y) \, dy$$

The first-order condition for this problem is:

$$\int y f_Y(y) \, dy = \int h f_Y(y) \, dy \implies h^* = E[Y]$$

Conditional Expectations as Orthogonal Projections

- Another perspective is that we can interpret the conditional expectation of Y given X as the orthgonal projection of Y onto the space of functions of the random variable X, i.e., L^2 space.

- This is the focus of the first two-ish weeks of Econ 2120.

- It provides a unifying perspective on much of regression analysis, and this is really the central focus of the first half of your econometrics sequence. It is an important idea that you'll spend a couple weeks thinking about.

Law of Iterated Expectations

- The law of iterated expectations is a really, really useful law for manipulating conditional expectations. It will show up all the time in your homework in a variety of settings.

- One form of the law of iterated expectations can be stated as:

$$E_Y[Y] = E_X E_{Y|X}[Y]$$

where E_X denotes the expectation taken with respect to the marginal density of X and $E_{Y|X}$ denotes the expectation taken with respect to the conditional density of Y given X.

- The way I think about this rule is as follows. Suppose I have the expectation of the conditional density of Y given X. If I take this expression and apply an expectation operator with respect to X, then I am left with the expectation of the marginal density of Y.

Break-Out Session #2

- We will now split out into breakout rooms to work through a couple additional problems.

- Please see the metrics math camp worksheet, available on GitHub:
  github.com/mdroste/metrics-mathcamp-2020

- Take 15 minutes to work through these problems with your small group. We'll re-convene to discuss any issues you might have had afterwards.

Day 1 Wrap-Up

- We covered a great deal of content today - you all earned a break!

- We will re-convene Monday to cover a little bit of additional information, and then we'll spend a much greater amount of time on exercises - particularly Matlab. So be sure to bring your laptop Monday, and have R/Matlab/Python installed (your choice).

- I will stick around for questions!